

Detecting Partial Data Writes in Delta Lake on Databricks

Methods Summary Table

Method	Detects Partial Writes?	Use Case
1. Delta DESCRIBE HISTORY	✓	Audit-based detection
2. Write Audit Logging	✓	Custom ETL pipelines
3. Streaming Checkpoints	✓	Structured Streaming
4. Row-level Hash Check	✓	Critical data integrity
5. Row Count Validation	✓	Batch and streaming
6. Delta Constraints	⚠ (limited)	Data quality enforcement
7. DLT Expectations	✓	Declarative pipeline QA

Sample.csv file

id, name, age
1, Alice, 30
2, Bob, 25
3, Charlie, 35
4, David, 40
5, Eva, 28

Step 0: Setup - Load the CSV and Write to Delta Table

```
python
CopyEdit
# Load sample CSV
df = spark.read.csv("/path/sample_data.csv", header=True, inferSchema=True)

# Write to Delta table
df.write.format("delta").mode("overwrite").save("/mnt/delta/sample_table")
```

1. Delta DESCRIBE HISTORY (Audit-based detection)

```
sql
CopyEdit
-- SQL version (in a notebook cell)
DESCRIBE HISTORY delta.`/mnt/delta/sample_table`;
Look at:


- operation: should be WRITE

```

Databricks: **Partial Write Patrol**

- operationMetrics.numOutputRows: should be 5
- userMetadata: optionally track job_id

If you see fewer rows than expected, that indicates a partial write.

It returns a history of all operations (like WRITE, MERGE, DELETE, etc.) on the Delta table, including:

Delta table, including:

version	timestamp	operation	operationMetrics	userMetadata
0	2025-07-09	WRITE	{"numOutputRows":"5"}	run_id=...
1	2025-07-09	WRITE	{"numOutputRows":"2"} ← suspicious	

2. Write Audit Logging (Custom ETL pipelines)

Field Descriptions:

Field	Type	Description
run_id	string	Unique ID for the ETL job run (generated with uuid)
table_name	string	Target Delta table name
expected_rows	integer	Number of rows you <i>expected</i> to be written
actual_rows	integer	Number of rows actually written (validated after the write)
write_success	boolean	True if the row counts match, else False
logged_at	string	Timestamp of the logging event

Add audit logs for tracking write activity:

python

CopyEdit

```
from datetime import datetime
```

```
import uuid
```

```
run_id = str(uuid.uuid4())
```

```
expected_rows = 5
```

```
actual_rows = spark.read.format("delta").load("/mnt/delta/sample_table").count()
```

```
log_df = spark.createDataFrame([
```

```
    "run_id": run_id,
```

```
    "table_name": "sample_table",
```

```
    "expected_rows": expected_rows,
```

```
    "actual_rows": actual_rows,
```

```
    "write_success": actual_rows == expected_rows,
```

```
    "logged_at": datetime.now().isoformat()
```

```
])
```

```
log_df.write.mode("append").saveAsTable("monitoring.write_audit_log")
```

run_id	table_name	expected_rows	actual_rows	write_success	logged_at
1c2b8a8e-1234-4cfd-9b00-abcde1234567	sample_table	5	3	False	2025-07-09T18:46:30.123456

append this above log to a monitoring table

#Later, query this audit table to find:

- Failed or partial writes
- Mismatched row counts
- Run-level tracing by run_id

3. Streaming with Checkpoints (Structured Streaming)

Step 1: Prepare Directory

Place your CSV (e.g., sample_data.csv) into a directory like:

/mnt/stream_input/

Databricks will treat each new file as a "mini-batch" of a stream

Step 2: Example Code - Structured Streaming from CSV to Delta

```
from pyspark.sql.functions import *
```

```
from pyspark.sql.types import StructType, IntegerType, StringType
```

```
# Define schema explicitly (recommended in streaming)
```

```
schema = StructType()
```

```
    .add("id", IntegerType())
```

```
    .add("name", StringType())
```

```
    .add("age", IntegerType())
```

```
# Read as a streaming DataFrame from the directory
```

```
stream_df = spark.readStream
```

```
    .schema(schema)
```

```
    .option("maxFilesPerTrigger", 1)  # "Only read 1 new file at a time (per trigger interval).
```

```
    .csv("/mnt/stream_input") # new CSVs added here will be streamed
```

```
# Write streaming data to Delta table with checkpointing
```

```
query = stream_df.writeStream
```

```
    .format("delta")
```

```
    .option("checkpointLocation", "/mnt/checkpoints/sample_streaming_table")
```

```
    .outputMode("append")
```

```
.start("/mnt/delta/sample_streaming_table")
```

How This Helps Detect Partial Writes

1) Atomicity and Checkpointing

- If a micro-batch fails mid-write, it does not commit to the Delta table.
- On recovery, the job retries from the last safe checkpoint.

2) Partial File Detection

- If a malformed CSV is written to /mnt/stream_input, it will fail parsing.
- You can capture those errors in the query listener or logs.

Output:

- Delta table: /mnt/delta/sample_streaming_table
- Checkpoints: /mnt/checkpoints/sample_streaming_table
- Triggered once per file (max 1 file per trigger due to maxFilesPerTrigger)

4. Row-level Hash Check (Critical data integrity)

python

CopyEdit

```
from pyspark.sql.functions import md5, concat_ws
```

```
df = df.withColumn("row_hash", md5(concat_ws("||", *df.columns)))
```

```
# Write with hash
```

```
df.write.format("delta").mode("overwrite").save("/mnt/delta/hashed_table")
```

```
# Validate written data
```

```
df_written = spark.read.format("delta").load("/mnt/delta/hashed_table")
```

```
# Detect any row differences
```

```
discrepancy = df.exceptAll(df_written)
```

```
if discrepancy.count() > 0:
```

```
    print("Partial write or corruption detected")
```

id	name	age	row_hash
1	Alice	30	0b3c0c8e32b4a4ddfca9c63c7d83b515

Example Result Table Format

Let's say this was your original data (df):

id	name	age	row_hash
1	Alice	30	34b7da764b21d298ef307d04d8152dc5
2	Bob	25	4e07408562bedb8b60ce05c1decfe3ad

3	Charlie	35	e1671797c52e15f763380b45e841ec32
---	---------	----	----------------------------------

But df_written (read back from Delta) only has:

id	name	age	row_hash
1	Alice	30	34b7da764b21d298ef307d04d8152dc5
2	Bob	25	4e07408562bedb8b60ce05c1decfe3ad

Then discrepancy.count() would return 1, and the discrepancy table would be:

id	name	age	row_hash
3	Charlie	35	e1671797c52e15f763380b45e841ec32

Summary Table

Component	Value
Total rows in df	3
Total rows in df_written	2
Discrepancy rows	1 (missing Charlie)
.count() > 0 result	☑ True
Message printed	"Partial write or corruption detected"

5. Row Count Validation (Batch and streaming)

python

CopyEdit

```
expected_count = 5
```

```
actual_count = spark.read.format("delta").load("/mnt/delta/sample_table").count()
```

```
if actual_count != expected_count:
```

```
    raise Exception(f"Partial write detected! Expected {expected_count}, got {actual_count}")
```

Example Scenarios:		Case 2: Partial Write Detected	
☑ Case 1: Successful Write			
Condition	Value	Condition	Value
expected_count	5	expected_count	5
actual_count	5	actual_count	3
Condition Met?	✗ (No Exception)	Condition Met?	☑ (Mismatch)
Result	☑ No error, write is valid	Result	✗ Raises Exception

Output :

Traceback (most recent call last):

...

Exception: Partial write detected! Expected 5, got 3

This clearly tells you that the Delta table is missing rows — only 3 out of 5 were written

To avoid hardcoding expected row counts, you can:

- Compare with `df.count()` before the write
- Track expected counts in a log or metadata store

6. Delta Constraints (Data quality enforcement)

Delta constraints help indirectly by catching missing/invalid values:

sql

CopyEdit

– In SQL notebook cell

```
ALTER TABLE delta.`/mnt/delta/sample_table`
```

```
ADD CONSTRAINT id_not_null CHECK (id IS NOT NULL);
```

Now if you write rows with null IDs, the write will fail and trigger alerts.

Adds a data constraint to the Delta table:

- Enforces that every row in the table must have a non-null id.
- Prevents writing invalid data into the table.
- Works like a guardrail to catch issues at write time.

What Happens Internally

Step 1: Constraint added

The constraint becomes metadata in the Delta transaction log. You can check it using:

sql

CopyEdit

```
SHOW TBLPROPERTIES delta.`/mnt/delta/sample_table`
```

Test Case: Try to Write a Bad Row

```
from pyspark.sql import Row
```

```
# Create a row with a null id
```

```
bad_df = spark.createDataFrame([Row(id=None, name="Test", age=33)])
```

```
# Try to append this invalid row to the constrained table
```

```
bad_df.write.format("delta").mode("append").save("/mnt/delta/sample_table")
```

Expected Output (Error)

If you run the write operation, Databricks throws an error like:

AnalysisException: CHECK constraint id_not_null (id IS NOT NULL) violated by row with values [null, Test, 33]

Field	Value
id	null

name	Test
age	33
Result	✗ Write Fails
Why	Constraint id IS NOT NULL is violated

When Write Succeeds
Summary Table

Scenario	Constraint Satisfied?	Write Outcome	Message
All rows have valid IDs	✓	✓ Success	No error
Some rows have id = NULL	✗	✗ Fail	CHECK constraint violated

Best Practice:

Use constraints to:

- Enforce data quality at write time
- Prevent accidental bad writes
- Pair with expectations in DLT for auto-monitoring

Would you like to see how to list all constraints on a Delt

7. Delta Live Tables (DLT) Expectations (Declarative QA)

In a DLT pipeline (Python syntax), use expectations:

python
CopyEdit

```
@dlt.table
@dlt.expect("valid_id", "id IS NOT NULL")
def cleaned_data():
    return spark.read.format("csv").option("header", True).load("/path/sample_data.csv")
```

DLT will automatically log failed rows and job status, preventing bad writes.

This is a DLT pipeline step that:

1. Creates a managed Delta Live Table called cleaned_data
2. Applies a data quality expectation that the id column must not be null
3. Automatically tracks, enforces, and logs this expectation in Databricks

@dlt.table

- Declares this function as a DLT table.
- The return value of the function will be materialized as a Delta table.
- Equivalent to writing a SQL CREATE TABLE AS SELECT.

Databricks: **Partial Write Patrol**

`@dlt.expect("valid_id", "id IS NOT NULL")`

- Adds a data quality rule (expectation) named "valid_id".
- Condition: id IS NOT NULL
- DLT tracks how many rows pass/fail this rule.

✓ Passed rows go into the table

✗ Failed rows are logged and dropped (or redirected)

`def cleaned_data():`

- Defines the logic for generating the table named cleaned_data.

`return spark.read.format("csv")...`

- Reads a CSV file as a Spark DataFrame and returns it for DLT processing.

Behind the Scenes – DLT Monitoring

DLT adds built-in observability:

Feature	What It Does
Quality dashboard	Shows how many records passed/failed each expectation
Auto logging	Logs metadata, errors, row counts, job run ID
Row dropping	Automatically drops rows that fail expectations
Error redirection	(Optional) Redirect bad rows to a quarantine table

Appendix:

@step 2:

Why It's Useful : maxFilesPerTrigger", 1

Benefit	Description
Simulate real-time ingestion	Useful in dev/test to mimic streaming from slowly-arriving files
Avoid processing overload	Prevents reading too many files in one batch, which can cause memory issues
Control ingestion rate	Helps apply backpressure and avoid spikes in processing

Example Use Case Assume your input directory /mnt/stream_input/ has these files: CopyEdit sample_data_1.csv sample_data_2.csv sample_data_3.csv With maxFilesPerTrigger = 1: <ul style="list-style-type: none"> • Micro-batch 1 → reads sample_data_1.csv • Micro-batch 2 → reads sample_data_2.csv • Micro-batch 3 → reads sample_data_3.csv Instead of reading all files at once, Spark spaces them out—1 file per micro-batch.	Without This Option If you don't set it, Spark may read all available files in the directory in a single micro-batch, which may lead to: <ul style="list-style-type: none"> • Heavy memory usage • Slower processing • Poor simulation of real-time data flow You can increase the value for higher throughput: .option("maxFilesPerTrigger", 5) # Process 5 files per trigger Or remove it for unlimited files per batch (default behavior).
---	--

@ Step 4:

```
df = df.withColumn("row_hash", md5(concat_ws("||", *df.columns)))
```

Purpose:

This line adds a new column called "row_hash" to your DataFrame (df) that contains a hash value for each row.

Explanation:

Expression Part	Meaning
*df.columns	Expands the list of column names to pass each column separately as an argument (e.g., "id", "name", "age")
`	
concat_ws(" md5(...)	Computes an MD5 hash of that concatenated string (e.g., `1
withColumn("row_hash", ...)	Adds a new column to the DataFrame called row_hash with the computed hash value

Why This Is Done:

To create a unique fingerprint of each row:

- Used for row-level integrity checking
- Helps detect partial writes, duplicate rows, or data corruption

Example:

Databricks: **Partial Write Patrol**

If the input row is:

id name age

1 Alice 30

The concatenated string will be:

"1| |Alice| |30"

Then the MD5 hash might be:

"0b3c0c8e32b4a4ddfca9c63c7d83b515"

So the final DataFrame becomes:

id	name	age	row_hash
1	Alice	30	0b3c0c8e32b4a4ddfca9c63c7d83b515

Use Case:

Later, after writing the data to a Delta table, you can re-read the table and recompute the row_hash again to compare with your source. If any hashes are missing or mismatched → the row was:

- Truncated
- Partially written
- Corrupted